

CONTENTS

EXERCISE A: NODE LEVEL STATISTICS	2
EXERCISE B: RANDOM WALKS	3
EXERCISE C: SHALLOW EMBEDDINGS	4
EXERCISE D: PROGRAMMING EXERCISE	5

Exercise A Node level statistics

Consider the following graph



The following code prints the adjacency matrix and computes its eigendecomposition:

```
>>> print(A)
[[0. 0. 1. 1. 0. 1. 0.]
 [0. 0. 0. 0. 1. 1. 1.]
 [1. 0. 0. 1. 0. 1. 0.]
 [1. 0. 1. 0. 0. 1. 0.]
 [0. 1. 0. 0. 0. 1. 1.]
 [1. 1. 1. 1. 1. 0. 1.]
 [0. 1. 0. 0. 1. 1. 0.]]

>>> lambda, E = np.linalg.eig(A)

>>> print(np.round(lambda, 3))
[ 3.646  2.   -1.646 -1.   -1.   -1.   -1.   ]

>>> print(np.round(E, 3))
[[-0.339 -0.408 -0.228 -0.816  0.004 -0.084  0.126]
 [-0.339  0.408 -0.228 -0.   -0.374  0.69  -0.255]
 [-0.339 -0.408 -0.228  0.408  0.511  0.114 -0.493]
 [-0.339 -0.408 -0.228  0.408 -0.515 -0.03  0.367]
 [-0.339  0.408 -0.228  0.   -0.176 -0.709 -0.377]
 [-0.558 -0.   0.83  -0.   0.   -0.   -0.   ]
 [-0.339  0.408 -0.228  0.   0.55  0.019  0.632]]
```

Question A.1: Determine the eigenvector centrality for each node in the graph.

Question A.2: Determine the clustering coefficient for each node in the graph.

Question B.1: Given a graph with adjacency matrix \mathbf{A} and a starting node chosen randomly according to a discrete distribution \mathbf{p} , what is the final node's probability distribution after taking a single step from the starting node along an edge chosen uniformly at random?

Question B.2: Given a graph with adjacency matrix \mathbf{A} , how many distinct paths of length t can we find starting from a specific node (say node 1)?

Hint: We can represent the initial state as a vector that is one for the start node and zero elsewhere. Where can we end up after taking a single step, and how can this be computed using a matrix-vector product? How can this approach be generalized to t steps?

In this exercise we will use the decoder $\text{DEC}(\mathbf{z}_u, \mathbf{z}_v) = \sigma(\mathbf{z}_u^\top \mathbf{z}_v + b)$, where $\sigma(x) = \frac{1}{1+e^{-x}}$ denotes the sigmoid function.

Question C.1: As a warm-up, show that $1 - \sigma(x) = \sigma(-x)$.

Based on this, can you spot a typo in eq. 3.12 on page 27 in the book?

Let $S_{u,v} \in \{0, 1\}$ denote a binary feature corresponding to an non-edge/edge between nodes u and v . Let $P_{u,v} = P(S_{u,v} = 1 | \mathbf{z}_u, \mathbf{z}_v, b) = \sigma(\mathbf{z}_u^\top \mathbf{z}_v + b)$ denote the predicted probability that the edge is present, given the latent node embeddings $\mathbf{z}_u, \mathbf{z}_v$ and bias b .

Question C.2: Write the cross entropy loss for the single observation $S_{u,v}$ (in terms of $P_{u,v}$ and $S_{u,v}$).

Question C.3: Let us assume that the embeddings \mathbf{z}_u and \mathbf{z}_v are orthogonal, such that their dot product is zero, and let us further assume that the bias is zero, $b = 0$. What is the probability of an edge between node u and v ?

Exercise D Programming exercise

In this exercise you will work with a shallow node embedding implemented in the script `shallow_embedding.py`. The code loads a graph from a file: This graph is simulated from a shallow embedding model, so that we know the ground truth probability of each possible link. In this exercise we will fit a shallow embedding model to the data and see how well we can estimate the ground truth.

Question D.1: Examine and run the code for loading the graph data.

- Understand how the graph is represented as a matrix as well as in the form of a set of index pairs and target values.
- It can perhaps help to visualize the adjacency matrix.

Question D.2: Examine and run the implementation of the class `Shallow`.

- Understand how the node embeddings are implemented using `torch.nn.Embedding`. Look up the documentation if needed.
- Understand what the forward function computes. What exactly is the role of the variables `rx` and `tx`?

Question D.3: Examine and run the code to fit the model. In this version, the loss is computed on the entire graph (no train/validation split and no mini batching).

- Experiment with different number of `max_step`.
- Experiment with different embedding dimensions. How does the embedding dimension influence the training loss?

Question D.4: Modify the code to use a train/validation split.

- Make a random split of the data (each node pair) into a training set (e.g. 80%) and a validation set (e.g. 20%).
- Modify the code to train on only the training data.
- Write code to compute the loss of the trained model on the validation set.
- Experiment with different embedding dimensions. What is the optimal embedding dimension when computing the loss on the validation set?

Question D.5: Hand in your predictions:

- Using the train/validation procedure you have implemented (or any other updates, hacks and modifications) to optimize the model. Compute what you believe is the best possible predicted link probability.
- Using the provided code, save your predictions in a file, `link_probabilities.pt`, and hand it in on DTU Learn.

I will compute the ground truth loss on your predictions and lowest generalization loss will be honored as the class winner.